



PUBLIC ANONYMIZED REPORT

The BridgeWorks Small-Business AI Visibility Report 2026

Why capable businesses remain difficult for AI-assisted search to understand

Public anonymized edition: 2026-07-16

Prepared by: BridgeWorks, Budapest

Evidence base: six scored audits, three public mini-scans, and one implementation verification

Primary audience: small-business owners across Africa and Central Europe

Executive summary

Small businesses are increasingly evaluated through search experiences that summarize, compare, and recommend options before a buyer visits a website. That includes traditional search, AI-generated summaries, ChatGPT search, Bing Copilot, and research tools that retrieve public web pages.

The practical question for a business is not whether it has added the phrase AI - ready to its marketing. The question is whether public systems can reliably discover and interpret accurate evidence about:

- who the business is;
- what it offers;
- where it operates;
- who is responsible for the expertise;
- why its claims should be trusted;
- which pages answer a buyer's question;
- what the buyer should do next.

BridgeWorks reviewed ten locally substantiated records for this first edition: six scored website audits, three narrower public mini-scans, and one implementation verification. The sample is small and non-random. It should be read as a field report from BridgeWorks' work, not as a market-wide prevalence study.

Within the six scored audits:

- 6 of 6 began with missing, near-absent, or materially weak structured business information.
- 6 of 6 had no `llms.txt` at the audited baseline. This is reported as an observed implementation gap, not as proof of a ranking disadvantage.
- 5 of 6 had a meaningful sitemap or URL-discovery problem, such as absent coverage, missing language versions, incomplete pages, or test URLs in production discovery files.
- 6 of 6 had a machine-readable entity gap involving the organization, named expert, location, or connection between first-party and third-party profiles.
- 5 of 6 had weak, thin, stale, or overly generic service evidence. One site had substantial niche content but failed to connect it to authorship, location, freshness, and structured identity.
- 6 of 6 had an absent, incomplete, unverified, or unobserved local-profile signal relevant to the business.

The three public mini-scans examined businesses for which no independent owned domain was located. In all three cases, reputation or marketplace activity existed, but treatment, service, team, location, and enquiry information remained fragmented across third-party platforms.

The central finding is simple:

Small-business AI visibility is usually not one missing technical file. It is a coordination problem across crawl access, entity clarity, content, structured information, third-party corroboration, and conversion.

The strongest first move is therefore not mass-producing generic AI content. It is creating one accurate, crawlable, well-evidenced source of truth for the business and connecting that source to the profiles, proof, and enquiry paths buyers already use.

1. What AI visibility means

BridgeWorks defines AI visibility as a business's readiness to be discovered, understood, corroborated, and cited across search and AI-assisted discovery surfaces.

That definition deliberately uses *readiness*. No agency controls the final output of a search engine or model. Results vary by query, geography, account context, freshness, available sources, competition, and platform behaviour.

AI visibility is also not a single official score. Google, OpenAI, Microsoft, and other platforms do not issue a shared 100-point business visibility rating. A BridgeWorks diagnostic score is an internal prioritisation tool. It helps an owner see which parts of the system are missing and whether the foundation improves after implementation.

Visibility is a chain

A business can fail at several different points:

1. Access: the page does not load, is blocked, or is hidden behind rendering or authentication problems.
2. Discovery: the page is not linked, not in a sitemap, or not known to the relevant index.
3. Understanding: the page is available, but the business identity, service, author, location, or relationship between facts is unclear.
4. Corroboration: the site makes claims that no reliable external profile, review, registry, publication, or customer evidence supports.
5. Selection: the source is understood, but other sources are more relevant, current, specific, or authoritative for the query.
6. Conversion: the business is discovered, but the buyer cannot confidently book, enquire, call, order, or compare the offer.

A technical audit that checks only schema can miss content and trust problems. A content audit that ignores crawler access can recommend articles that remain undiscovered. A profile audit that ignores the owned website can leave the business dependent on platforms it does not control.

2. What the platforms actually say

AI visibility is an emerging field, and its language is often more confident than the evidence. This report uses primary-source guidance wherever a platform has published it.

OpenAI search access is distinct from model training

OpenAI documents separate user agents for separate purposes. OAI-SearchBot is used for ChatGPT search results. GPTBot is associated with content that may be used for training foundation models. ChatGPT-User may visit a page following a user action and is not the crawler that determines Search inclusion.

For a business that wants to be eligible for ChatGPT search, the relevant technical check is whether OAI-SearchBot can access the public pages and whether the site or hosting layer blocks OpenAI's published IP ranges. Allowing GPTBot is a separate policy decision.

Source: <https://developers.openai.com/api/docs/bots>

Google says the fundamentals still apply

Google's guidance for AI Overviews and AI Mode says there are no additional technical requirements and no special AI schema needed beyond established Search fundamentals. Pages must still be accessible, indexable, useful, and eligible to appear in Search.

This matters because AI optimization should not be sold as a secret markup layer. Strong technical SEO, clear text, accurate structured data, useful pages, good internal discovery, and reliable business information remain the foundation.

Source: <https://developers.google.com/search/docs/appearance/ai-features>

Structured data helps understanding, not guaranteed inclusion

Google describes structured data as explicit clues about the meaning of a page. Accurate markup can help a search system understand organizations, people, products, services, articles, locations, and other entities.

Google also states that valid structured data does not guarantee a rich result. Markup must match visible page content and follow quality guidelines. BridgeWorks applies the same guardrail to AI visibility: schema is useful when it accurately describes a real, visible business. It is not a shortcut around weak content or unsupported claims.

Sources:

- <https://developers.google.com/search/docs/appearance/structured-data/intro-structured-data>
- <https://developers.google.com/search/docs/appearance/structured-data/sd-policies>

Bing emphasizes discovery and freshness

Bing recommends crawlable internal links, XML sitemaps, accurate canonical URLs, and IndexNow notifications. Its webmaster guidance connects those practices with accurate indexing and eligibility across Bing and Copilot experiences.

Sitemaps and IndexNow still do not guarantee ranking or selection. They reduce avoidable discovery and freshness failures.

Sources:

- <https://www.bing.com/webmasters/help/webmaster-guidelines-30fba23a>
- <https://www.bing.com/webmasters/help/sitemaps-3b5cf6ed>
- <https://www.bing.com/webmasters/help/indexnow-0z209wby>

llms.txt is experimental

The llms.txt project describes itself as a proposal for a concise Markdown overview that helps language models use a website at inference time. It is not a replacement for robots.txt or sitemap.xml, and it is not a confirmed ranking signal adopted by every major platform.

BridgeWorks may implement llms.txt as a low-cost, human-readable index of important pages. We do not score its absence as proof that a site cannot appear in AI answers, and we do not promise that adding it will create citations.

Source: <https://llmstxt.org/>

3. The six-part visibility system

Starting with methodology v1, BridgeWorks evaluates six connected domains.

1. Crawl and index readiness

The first question is whether public systems can access and discover the right URLs.

Checks include:

- reliable HTTP responses;
- redirect and canonical consistency;
- robots and page-level index controls;
- complete XML sitemaps;
- crawlable internal links;
- freshness signals such as accurate modification dates;
- Search Console or Bing evidence where access exists;
- relevant crawler access, including OAI-SearchBot when ChatGPT search eligibility matters.

This layer prevents the most basic failure: useful pages that cannot enter the candidate set.

2. Entity clarity

An entity is the identifiable business, person, place, or organization behind the pages. Small businesses often assume a logo and a contact page are enough. In practice, different sources may use different names, addresses, categories, phone numbers, or founder identities.

Checks include:

- consistent business name and location;
- legal and trading identity where appropriate;
- named founders, clinicians, advisors, architects, or directors;
- specific business category and service area;
- authoritative social and local profiles;
- clear relationships between the organization and its people;
- disambiguation from similarly named businesses.

An entity gap is especially damaging for founder-led professional services, where the person's qualifications are part of the buyer's decision.

3. Answer-ready content

AI-assisted systems retrieve and summarize sources that answer a question. Generic slogans provide very little evidence.

Answer-ready content includes:

- a specific description of each service;

- who the service is for and where it is available;
- process, constraints, pricing logic, or eligibility where publishable;
- named expertise and first-hand experience;
- dated examples, case evidence, policies, and outcomes;
- direct answers to buyer questions;
- clear authorship and update dates where freshness matters;
- citations to authoritative external sources for factual claims.

There is no universal minimum word count. A concise contact answer can be sufficient. A medical, financial, legal, or technical guide may require much more depth and stronger sourcing.

4. Structured information

Structured information labels facts in a machine-readable format. Relevant types may include Organization, LocalBusiness, ProfessionalService, Person, Article, Course, Product, BreadcrumbList, or other supported vocabulary.

Good implementation follows three rules:

1. It describes content that users can also see.
2. It uses accurate, specific properties rather than filling every possible field.
3. It is validated and monitored after deployment.

Structured data supports understanding and search features. It does not replace content, authority, or indexing.

5. Authority and corroboration

First-party pages explain what the business says about itself. Third-party evidence helps another system decide whether those claims are credible.

Useful corroboration may include:

- verified Google Business Profile details;
- genuine customer reviews;
- relevant directories and professional registries;
- memberships and qualifications;
- press and interviews;
- supplier or partner references;
- public project or client evidence;
- consistent social profiles.

The goal is not to create profiles everywhere. It is to make the important sources accurate, consistent, and connected.

6. Conversion and measurement

Visibility without a next step is not an acquisition system.

Checks include:

- working enquiry and booking forms;
- clear phone, email, WhatsApp, order, or scheduling actions;
- route-specific calls to action;
- analytics and search measurement;
- CRM or lead-routing evidence;
- consent and double opt-in where required;
- handover ownership and operating instructions.

This is where a visibility project becomes commercial. The final measure is not the audit score. It is qualified enquiries, conversations, bookings, proposals, and revenue.

4. Findings from the substantiated corpus

Finding 1: the business exists, but the entity does not resolve cleanly

All six scored audits contained a material machine-readable identity gap.

The pattern appeared in different forms:

- an HR consultancy with no visible founder or current named team;
- an architecture firm whose founder credentials existed on one long page but were not connected through structured information;
- a financial advisor with strong credentials and 25 years in Poland, but no structured location identity;
- a consumer brand whose founder story was central to the product but lacked a dedicated credential page and Person relationship;
- an education business with a practitioner founder and nine courses, but no organization, person, or course structure;
- a facility-management company whose offline history was not represented clearly online at baseline.

The commercial consequence is not merely technical. A buyer cannot confidently compare a business if the site does not make its people, location, category, and qualifications explicit.

Finding 2: schema was usually absent where the strongest facts already existed

All six scored audits began with missing, near-absent, or materially weak structured information.

In several cases, the underlying content was already valuable:

- credentials were written in an About page;
- testimonials were visible but not connected to the relevant person or service;
- detailed articles existed without structured author or date information;
- service pages existed without service or organization identity;
- courses existed without course and educator structure.

The opportunity was not to invent more claims. It was to label accurate, visible facts consistently.

Finding 3: URL discovery failed in ordinary ways

Five of six scored audits had a material sitemap or discovery-control problem.

Examples included:

- missing language versions;
- an absent sitemap;
- only a fraction of important pages listed;
- test URLs left in production discovery files;
- outdated modification dates;
- important new service pages omitted.

These are ordinary technical SEO failures, but they matter to AI-assisted search for the same reason they matter to traditional search: a system cannot select a page it has not discovered or indexed.

Finding 4: content depth and business proof were disconnected

Five scored audits had weak, stale, thin, or generic service evidence.

One architecture site had a strong 30-plus-year story but project pages that rarely explained the challenge, work, or result. One HR site named real clients but hid current leadership and relied on old, generic copy. One education site had real practitioner expertise but no durable glossary, student outcome evidence, or structured course authority.

The exception was the financial-advice site. It had detailed niche guides and genuine credentials. Its weakness was connection: missing author/date structure, weak local identity, absent service schema, and limited linking between the Poland expertise and the commercial offer.

This distinction matters. Not every business needs more content. Some need better packaging and evidence architecture around content they already have.

Finding 5: local trust was treated as a separate platform

All six scored audits had a local-profile signal that was absent, incomplete, unverified, or not observable in the audit.

For location-dependent services, the website and local profile should corroborate each other:

- business name;
- category;
- address or service area;
- phone and opening hours;
- services;
- reviews;
- website and booking route.

When these disagree, a person has to resolve the conflict manually. Machines face the same ambiguity.

Finding 6: some businesses do not own the destination

No independent owned domain was located for any of the three mini-scan businesses.

Each still had public proof:

- a Lagos dental clinic had directory and Maps visibility;
- a Budapest dental/aesthetic clinic had award and directory references;
- a Warsaw rental agency had an active marketplace identity and business record.

The problem was control. None had one located source where service details, team identity, buyer questions, trust proof, language support, and direct enquiry could accumulate together.

Marketplaces and directories can remain valuable. An owned site does not replace them. It gives the business a canonical source that those platforms can point toward.

5. An anonymized B2B services engagement

For an established Central European facilities-services company, BridgeWorks recorded a GEO and AI-visibility diagnostic score of 16/100 at baseline.

The work connected:

- a bilingual server-rendered website;
- clearer service and contact routes;
- structured business information;
- legal and entity consistency;
- sitemap and crawler controls;
- Bing Webmaster and IndexNow;
- an experimental `llms.txt` overview;
- content and reporting;
- a targeted prospect database and outreach sequence;
- a documented handover package.

The engagement-close audit recorded 77/100, a 61-point improvement from the baseline.

That score movement is evidence that the audited infrastructure and content signals improved. It is not evidence that the client appears for every relevant prompt, nor is it a revenue attribution claim.

The close audit still listed unresolved items: Google Business Profile verification, sitemap coverage for four new funnel pages, FAQPage markup, named-team depth, and third-party review evidence. Those limitations are part of the case, not footnotes to hide.

The commercial lesson is that the client needed coordination before amplification. Advertising a business whose service pages, identity, discovery, and enquiry systems are unclear spends money on top of ambiguity.

6. Why incomplete proof stays private

BridgeWorks also reviewed an active hospitality implementation. It is included only as an anonymized process observation, not as a completed outcome case study.

The new preview includes:

- a clear restaurant identity and location;
- structured menu data;
- menu, contact, catering, order, and WhatsApp routes;
- structured website information;
- an AI-readable business overview;
- a connected MailerLite group and welcome automation;
- responsive desktop and mobile evidence.

At the evidence cut-off, the approved root domain still served the legacy site. New sitemap and robots routes passed the local build but were not deployed. The retention path had not completed an approved end-to-end subscriber test. Search tooling, client cutover approval, final platform evidence, handover acceptance, and production metrics remained open.

That is why BridgeWorks does not use this engagement as finished performance proof. The correct sequence is:

1. finish the delivery;
2. verify the live system;
3. capture approval and limitations;
4. then extract the case study and content.

7. The practical small-business checklist

Identity

- The same public business name is used across the site and primary profiles.
- Location or service area is explicit.
- Business category and core offer are clear in one sentence.
- Founders, clinicians, advisors, architects, or other experts are named where relevant.
- Qualifications and memberships are accurate and verifiable.
- Similar names or locations are disambiguated.

Access and discovery

- Important public pages return reliable 200 responses.
- HTTPS, redirects, and canonical URLs are consistent.
- Robots rules reflect the owner's actual search and training preferences.
- OAI - SearchBot access is checked separately from GPTBot policy.
- The XML sitemap lists canonical, indexable, current URLs.
- Important pages are linked from crawlable navigation or related pages.
- Google Search Console and Bing Webmaster are connected where appropriate.
- Updated URLs are submitted or discoverable through normal platform tools.

Content

- Every core service has a dedicated, useful page.
- Pages answer who, what, where, process, constraints, and next-step questions.
- Claims use first-hand evidence or authoritative citations.
- Articles show authorship and dates when freshness matters.
- Case evidence includes context, work, result, and limitations.
- FAQs answer genuine buyer questions rather than keyword variations.

Structured information

- Organization or business markup matches visible facts.
- Relevant people, services, locations, articles, products, or courses use appropriate structure.
- Structured data is validated before and after deployment.
- Reviews, ratings, credentials, and prices are not invented or hidden from users.
- Entity relationships use stable URLs and identifiers where appropriate.

Authority and local proof

- The main local/business profile is claimed and accurate.
- Reviews are genuine, recent enough to be useful, and answered appropriately.
- Important directories and registries agree on identity and contact details.
- Professional profiles connect named experts to the business.
- Press, partner, membership, or project references are linked where permitted.

Conversion and measurement

- Forms, phone links, WhatsApp links, booking links, and order routes work on mobile.
- Every important page has one logical next step.
- Consent, double opt-in, and unsubscribe controls work where required.
- Search and analytics accounts are owned and documented.
- Enquiries reach an accountable owner or CRM stage.
- Before-and-after evidence is captured before making an outcome claim.

Experimental layer

- Consider a concise `llms.txt` overview as an experiment, not a ranking promise.
- Record dated platform query tests with exact prompts and citations.
- Re-test important prompts rather than treating one answer as a permanent ranking.

8. What to fix before spending more on ads

Ads can create traffic. They cannot make an unclear business easier to evaluate.

Before increasing spend, confirm five things:

1. The landing page matches the offer. A buyer should not have to infer which service the advertisement refers to.
2. The business is identifiable. Name, people, location, category, and contact details should be consistent.
3. Trust is visible. Relevant reviews, evidence, qualifications, policies, and limitations should be near the decision.
4. The next step works. Test the form, booking, call, WhatsApp, or checkout path on a real mobile device.
5. Measurement exists. Know which campaign and page produced the enquiry and whether the enquiry became a qualified conversation.

For many small businesses, fixing those foundations improves both human conversion and machine understanding. That is a more defensible claim than promising an AI ranking.

9. A 30-day action plan

Week 1: establish the source of truth

- Reconcile business name, location, category, phone, email, and primary profiles.
- Identify the five to ten pages that should explain the business.
- Fix access, redirect, canonical, robots, and sitemap errors.
- Connect Search Console and Bing Webmaster.

Week 2: structure the business evidence

- Rewrite the core business and service descriptions in direct language.
- Add named experts, qualifications, dates, and visible proof.
- Implement accurate structured information for the organization and highest-value pages.
- Validate the deployed pages.

Week 3: build answer-ready assets

- Publish one substantial page answering a high-value buyer question.
- Add genuine FAQs to the relevant service pages.
- Turn one project or customer outcome into a transparent case example.
- Improve internal links between authority content and commercial pages.

Week 4: connect discovery to conversion

- Reconcile local and professional profiles.
- Test all enquiry and booking routes.
- Add source measurement and CRM ownership.
- Run a small set of dated platform query tests.
- Record the baseline and decide which gaps deserve the next month of work.

10. How to use an AI Visibility Scan

A useful scan should answer three questions:

1. What evidence can public systems discover today?
2. Where does the chain from discovery to understanding to enquiry break?
3. Which three fixes have the highest commercial value for this business?

It should not be a list of every possible schema type or a promise of guaranteed AI citations. The output should distinguish observed facts, reported inputs, inferences, and untested areas. It should end with a small, prioritized implementation plan.

Conclusion

Small businesses do not need to become publishers of generic AI news. They need to become clear, trustworthy sources about their own work.

That means:

- accessible pages;
- a resolvable business entity;
- specific answers and named expertise;
- accurate structured information;
- corroborating profiles and proof;
- a working path to enquiry;
- measurement that connects visibility to commercial outcomes.

The opportunity is not to manipulate an AI system. It is to remove the ambiguity that prevents buyers and machines from understanding a capable business.

Request a scan

BridgeWorks provides evidence-led AI Visibility Scans for small businesses across Africa and Central Europe.

The scan reviews the public website, entity signals, structured information, discovery controls, authority evidence, local profiles, and conversion path. It returns prioritized actions and clearly states what was observed, inferred, and not tested.

Contact: office@bridgeworks.agency

Method and evidence notes

- Full method: `operations/client-acquisition-engine/research/AI-VISIBILITY-METHODOLOGY-V1-2026-07-14.md`
- Anonymised evidence table: `operations/client-acquisition-engine/research/ai-visibility-evidence-dataset-2026-07-14.csv`
- Named client evidence: retained in the private BridgeWorks evidence register.
- Public case-study source: `public-proof/CASE-STUDY-CENTRAL-EUROPEAN-B2B-SERVICES-2026.md`
- Active hospitality verification: retained privately and not released as completed proof.